

# ML PROJECT 2 - VEHICLE PRICE PREDICTION

our algorithm will take vehicle details like mileage,engin type, no of doors, lenght,width,height,engine capacity,etc.... and our algorithm will predict PRICE of the vehicle

## step 1 - load the data

```
In [1]: 1 import pandas as pd
        2
        3 auto_data = pd.read_csv('auto.txt')
        4 auto_data.head()
```

Out[1]:

	symboling	normalized- losses	make	fuel- type	aspiration	num- of- doors	body- style	drive- wheels	engine- location	wheel- base	
0	3	?	alfa- romero	gas	std	two	convertible	rwd	front	88.6	.
1	3	?	alfa- romero	gas	std	two	convertible	rwd	front	88.6	.
2	1	?	alfa- romero	gas	std	two	hatchback	rwd	front	94.5	.
3	2	164	audi	gas	std	four	sedan	fwd	front	99.8	.
4	2	164	audi	gas	std	four	sedan	4wd	front	99.4	.

5 rows × 26 columns



## step 2 - clean data

```
In [2]: 1 # you can observe there are some columns with values ? Lets clean these.
        2 import numpy as np
        3 auto_data = auto_data.replace('?', np.nan)
        4 auto_data.head()
```

Out[2]:

	symboling	normalized- losses	make	fuel- type	aspiration	num- of- doors	body- style	drive- wheels	engine- location	wheel- base	.
0	3	NaN	alfa-romero	gas	std	two	convertible	rwd	front	88.6	.
1	3	NaN	alfa-romero	gas	std	two	convertible	rwd	front	88.6	.
2	1	NaN	alfa-romero	gas	std	two	hatchback	rwd	front	94.5	.
3	2	164	audi	gas	std	four	sedan	fwd	front	99.8	.
4	2	164	audi	gas	std	four	sedan	4wd	front	99.4	.

5 rows × 26 columns



```
In [3]: 1 auto_data['price'].describe() # Let see what is the data type of price column
```

```
Out[3]: count      201
        unique     186
        top       8921
        freq        2
        Name: price, dtype: object
```

```
In [4]: 1 # we have to convert this column to float type
        2 auto_data['price'] = pd.to_numeric(auto_data['price'], errors='coerce') #coe
        3 auto_data['price'].describe() #now type is converted to float
```

```
Out[4]: count      201.000000
        mean     13207.129353
        std       7947.066342
        min       5118.000000
        25%       7775.000000
        50%      10295.000000
        75%      16500.000000
        max      45400.000000
        Name: price, dtype: float64
```

```
In [5]: 1 # Let us remove unwanted columns -- which are not useful
```

```
In [6]: 1 auto_data = auto_data.drop('normalized-losses', axis=1)
        2 auto_data.head()
```

Out[6]:

	symboling	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	length	...	en
0	3	alfa-romero	gas	std	two	convertible	rwd	front	88.6	168.8	...	
1	3	alfa-romero	gas	std	two	convertible	rwd	front	88.6	168.8	...	
2	1	alfa-romero	gas	std	two	hatchback	rwd	front	94.5	171.2	...	
3	2	audi	gas	std	four	sedan	fwd	front	99.8	176.6	...	
4	2	audi	gas	std	four	sedan	4wd	front	99.4	176.6	...	

5 rows × 25 columns



```
In [7]: 1 auto_data.columns
```

Out[7]: Index(['symboling', 'make', 'fuel-type', 'aspiration', 'num-of-doors', 'body-style', 'drive-wheels', 'engine-location', 'wheel-base', 'length', 'width', 'height', 'curb-weight', 'engine-type', 'num-of-cylinders', 'engine-size', 'fuel-system', 'bore', 'stroke', 'compression-ratio', 'horsepower', 'peak-rpm', 'city-mpg', 'highway-mpg', 'price'], dtype='object')

```
In [8]: 1 auto_data['horsepower'].describe()
```

Out[8]: count 203  
unique 59  
top 68  
freq 19  
Name: horsepower, dtype: object

```
In [9]: 1 auto_data['horsepower'] = pd.to_numeric(auto_data['horsepower'], errors='coerce')
        2 auto_data['horsepower'].describe()
```

Out[9]: count 203.000000  
mean 104.256158  
std 39.714369  
min 48.000000  
25% 70.000000  
50% 95.000000  
75% 116.000000  
max 288.000000  
Name: horsepower, dtype: float64

```
In [10]: 1 auto_data.columns
```

```
Out[10]: Index(['symboling', 'make', 'fuel-type', 'aspiration', 'num-of-doors',  
              'body-style', 'drive-wheels', 'engine-location', 'wheel-base', 'length',  
              'width', 'height', 'curb-weight', 'engine-type', 'num-of-cylinders',  
              'engine-size', 'fuel-system', 'bore', 'stroke', 'compression-ratio',  
              'horsepower', 'peak-rpm', 'city-mpg', 'highway-mpg', 'price'],  
              dtype='object')
```

```
In [11]: 1 auto_data['bore']
```

```
Out[11]: 0      3.47  
         1      3.47  
         2      2.68  
         3      3.19  
         4      3.19  
         ...  
        200     3.78  
        201     3.78  
        202     3.58  
        203     3.01  
        204     3.78  
         Name: bore, Length: 205, dtype: object
```

```
In [12]: 1 auto_data['bore'] = pd.to_numeric(auto_data['bore'], errors='coerce')#cnvrt  
         2 auto_data['bore'].describe()
```

```
Out[12]: count      201.000000  
         mean         3.329751  
         std          0.273539  
         min          2.540000  
         25%          3.150000  
         50%          3.310000  
         75%          3.590000  
         max          3.940000  
         Name: bore, dtype: float64
```

```
In [13]: 1 auto_data['stroke'] = pd.to_numeric(auto_data['stroke'], errors='coerce')#cn  
         2 auto_data['stroke'].describe()
```

```
Out[13]: count      201.000000  
         mean         3.255423  
         std          0.316717  
         min          2.070000  
         25%          3.110000  
         50%          3.290000  
         75%          3.410000  
         max          4.170000  
         Name: stroke, dtype: float64
```

```
In [14]: 1 auto_data['peak-rpm'] = pd.to_numeric(auto_data['peak-rpm'], errors='coerce')
        2 auto_data['peak-rpm'].describe()
```

```
Out[14]: count      203.000000
         mean      5125.369458
         std       479.334560
         min      4150.000000
         25%      4800.000000
         50%      5200.000000
         75%      5500.000000
         max      6600.000000
         Name: peak-rpm, dtype: float64
```

```
In [15]: 1 auto_data['num-of-cylinders']
```

```
Out[15]: 0      four
         1      four
         2      six
         3      four
         4      five
         ...
        200     four
        201     four
        202     six
        203     six
        204     four
         Name: num-of-cylinders, Length: 205, dtype: object
```

```
In [16]: 1 cylinders_dict = {
2         'two':2,
3         'three':3,
4         'four':4,
5         'five':5,
6         'six':6,
7         'eight':8,
8         'twelve':12
9     }
10 auto_data['num-of-cylinders'].replace(cylinders_dict, inplace = True)
11 auto_data.head()
```

Out[16]:

	symboling	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	length	...
0	3	alfa-romero	gas	std	two	convertible	rwd	front	88.6	168.8	...
1	3	alfa-romero	gas	std	two	convertible	rwd	front	88.6	168.8	...
2	1	alfa-romero	gas	std	two	hatchback	rwd	front	94.5	171.2	...
3	2	audi	gas	std	four	sedan	fwd	front	99.8	176.6	...
4	2	audi	gas	std	four	sedan	4wd	front	99.4	176.6	...

5 rows × 25 columns

```
In [17]: 1 auto_data['num-of-cylinders'].head()
```

```
Out[17]: 0    4
1    4
2    6
3    4
4    5
Name: num-of-cylinders, dtype: int64
```

```

In [18]: 1 a = {'1bbl':1, '2bbl':2, '4bbl':4, 'idi':5, 'mfi':6, 'mpfi':7,
2         'spdi':8, 'spfi':9}
3 auto_data['fuel-system'].replace(a, inplace = True)
4
5 b = {'dohc':1, 'dohcv':2, 'l':3, 'ohc':4, 'ohcf':5, 'ohcv':6,
6      'rotor':7}
7 auto_data['engine-type'].replace(b, inplace = True)
8
9 c = {'front':1, 'rear':2}
10 auto_data['engine-location'].replace(c, inplace = True)
11
12 d = {'4wd':1, 'fwd':2, 'rwd':3}
13 auto_data['drive-wheels'].replace(d, inplace = True)
14
15 e = {
16     'alfa-romero' : 1,'audi' : 2,'bmw': 3,'chevrolet' :4,'dodge':5,
17     'honda':6, 'isuzu':7,'jaguar':8, 'mazda':9, 'mercedes-benz':10,
18     'mercury':11,'mitsubishi':12, 'nissan':13, 'peugot':14,
19     'plymouth':15,'porsche':16, 'renault':17, 'saab':18, 'subaru':19,
20     'toyota':20, 'volkswagen':21, 'volvo':22
21 }
22 auto_data['make'].replace(e, inplace = True)
23
24 f = {'convertible':1,'hardtop':2, 'hatchback':3, 'sedan':4, 'wagon':5}
25 auto_data['body-style'].replace(f, inplace = True)
26
27 g = {'four':4, 'two':2}
28 auto_data['num-of-doors'].replace(g, inplace = True)
29
30 h = {'std':0, 'turbo':1}
31 auto_data['aspiration'].replace(h, inplace = True)
32
33
34
35 auto_data.head(10)

```

Out[18]:

	symboling	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	length	...	eng
0	3	1	gas	0	2.0	1	3	1	88.6	168.8	...	
1	3	1	gas	0	2.0	1	3	1	88.6	168.8	...	
2	1	1	gas	0	2.0	3	3	1	94.5	171.2	...	
3	2	2	gas	0	4.0	4	2	1	99.8	176.6	...	
4	2	2	gas	0	4.0	4	1	1	99.4	176.6	...	
5	2	2	gas	0	2.0	4	2	1	99.8	177.3	...	
6	1	2	gas	0	4.0	4	2	1	105.8	192.7	...	
7	1	2	gas	0	4.0	5	2	1	105.8	192.7	...	
8	1	2	gas	1	4.0	4	2	1	105.8	192.7	...	
9	0	2	gas	1	2.0	3	1	1	99.5	178.2	...	

10 rows × 25 columns



```
In [19]: 1 i = {'gas':0, 'diesel':1}
          2 auto_data['fuel-type'].replace(i,inplace = True)
          3
          4 auto_data.head(10)
```

Out[19]:

	symboling	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	length	...	engine-size
0	3	1	0	0	2.0	1	3	1	88.6	168.8	...	130
1	3	1	0	0	2.0	1	3	1	88.6	168.8	...	130
2	1	1	0	0	2.0	3	3	1	94.5	171.2	...	152
3	2	2	0	0	4.0	4	2	1	99.8	176.6	...	109
4	2	2	0	0	4.0	4	1	1	99.4	176.6	...	136
5	2	2	0	0	2.0	4	2	1	99.8	177.3	...	136
6	1	2	0	0	4.0	4	2	1	105.8	192.7	...	136
7	1	2	0	0	4.0	5	2	1	105.8	192.7	...	136
8	1	2	0	1	4.0	4	2	1	105.8	192.7	...	131
9	0	2	0	1	2.0	3	1	1	99.5	178.2	...	131

10 rows × 25 columns





```
In [20]: 1 auto_data.isnull()
```

Out[20]:

	symboling	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	length	...	engir si
0	False	False	False	False	False	False	False	False	False	False	...	Fal
1	False	False	False	False	False	False	False	False	False	False	...	Fal
2	False	False	False	False	False	False	False	False	False	False	...	Fal
3	False	False	False	False	False	False	False	False	False	False	...	Fal
4	False	False	False	False	False	False	False	False	False	False	...	Fa
...	...	...	...	...	...	...	...	...	...	...	...	
200	False	False	False	False	False	False	False	False	False	False	...	Fal
201	False	False	False	False	False	False	False	False	False	False	...	Fal
202	False	False	False	False	False	False	False	False	False	False	...	Fal
203	False	False	False	False	False	False	False	False	False	False	...	Fal
204	False	False	False	False	False	False	False	False	False	False	...	Fal

205 rows × 25 columns



```
In [21]: 1 auto_data.isna().sum()
```

```
Out[21]: symboling      0
make      0
fuel-type 0
aspiration 0
num-of-doors 2
body-style 0
drive-wheels 0
engine-location 0
wheel-base 0
length 0
width 0
height 0
curb-weight 0
engine-type 0
num-of-cylinders 0
engine-size 0
fuel-system 0
bore 4
stroke 4
compression-ratio 0
horsepower 2
peak-rpm 2
city-mpg 0
highway-mpg 0
price 4
dtype: int64
```

```
In [22]: 1 # horsepower
2 # curb-weight
3 # peak-rpm
4 del auto_data['horsepower']
5 del auto_data['curb-weight']
6 del auto_data['peak-rpm']
7
```

```
In [ ]: 1
```

```
In [23]: 1 # Lets clean up our data
2 auto_data = auto_data.dropna()
```

## step 3 -train test split

```
In [24]: 1 from sklearn.model_selection import train_test_split
2         # lets feed our data to machine learning model
3         x = auto_data.drop('price', axis=1)
4
5         y = auto_data['price']
6
7         # split data into 80% for training, 20% for testing
8         x_train, x_test, y_train, y_test = train_test_split(x,y, test_size = 0.2, ra
```

## step 4 - train the algorithm

```
In [25]: 1 from sklearn.linear_model import LinearRegression
2
3         linear_model = LinearRegression()
4         linear_model.fit(x_train, y_train)
```

Out[25]: LinearRegression()

```
In [26]: 1 linear_model.score(x_train, y_train) # traing data accuracy
```

Out[26]: 0.8977925747050544

```
In [27]: 1 linear_model.coef_
```

Out[27]: array([ -667.83607138, -201.03840036, -11877.75594957, 3267.45664655,  
 -341.88511643, -275.67970494, 534.24818935, 12457.43964247,  
 -27.57274928, 13.49191389, 735.2578018 , 331.25941981,  
 -633.39734035, 167.85881525, 149.24314962, 266.08507362,  
 -3020.67818985, -4381.08882035, 866.25626477, -280.79536587,  
 314.52728571])

```
In [28]: 1 predictors = x_train.columns #see the weights associated with particular fe
2         coef = pd.Series(linear_model.coef_,predictors).sort_values()
3         print(coef)
```

```
fuel-type          -11877.755950
stroke             -4381.088820
bore               -3020.678190
symboling          -667.836071
engine-type        -633.397340
num-of-doors       -341.885116
city-mpg           -280.795366
body-style         -275.679705
make               -201.038400
wheel-base        -27.572749
length             13.491914
engine-size        149.243150
num-of-cylinders   167.858815
fuel-system        266.085074
highway-mpg        314.527286
height             331.259420
drive-wheels       534.248189
width              735.257802
compression-ratio  866.256265
aspiration         3267.456647
engine-location    12457.439642
dtype: float64
```

## step 5 - test the algorithm


```
In [29]: 1 # Lets predict using linear regression model
2         y_predict = linear_model.predict(x_test)
```

```
In [30]: 1 # Lets plot the prediction using matplotlib lib
2 %pylab inline
3 pylab.rcParams['figure.figsize'] = (15,6)
4
5 plt.plot(y_predict, label='Predicted')
6 plt.plot(y_test.values, label='Actual')
7
8 plt.ylabel('price')
9 plt.legend()
10 plt.show()
```

Populating the interactive namespace from numpy and matplotlib

C:\Users\91918\AppData\Local\Programs\Python\Python39\lib\site-packages\IPython\core\magics\pylab.py:159: UserWarning: pylab import has clobbered these variables: ['f', 'e']

`%matplotlib` prevents importing \* from pylab and numpy  
warn("pylab import has clobbered these variables: %s" % clobbered +



```
In [31]: 1 # how well our regression model works on our test data
2 score = linear_model.score(x_test, y_test)
3 score
```

Out[31]: 0.8409940243694667

**step 6 - find the error(how much is the error in the output)**

**MEAN SQUARED ERROR --- ### this will tell us how much error is there in the given output.**

```
In [32]: 1 from sklearn.metrics import mean_squared_error
2
3 linear_model_mse = mean_squared_error(y_predict, y_test) #predicted y and ac
4 linear_model_mse # its coming out to be 26 millions
```

Out[32]: 9848274.421919573

```
In [33]: 1 import math
2 math.sqrt(linear_model_mse)
```

Out[33]: 3138.1960458071408

**conclusion for linear regression algorithm - accuracy is 84% and error is 3,138 dollars**

## step 7 : improvie, let us test other algorithms

```
In [34]: 1 # implementing lasso and ridge regression models.
2 from sklearn.linear_model import Lasso
3
4 lasso_model = Lasso(alpha=0.5, normalize=True) #alpha is regularization para
5 lasso_model.fit(x_train, y_train)
```

C:\Users\91918\AppData\Local\Programs\Python\Python39\lib\site-packages\sklearn\linear\_model\\_base.py:141: FutureWarning: 'normalize' was deprecated in version 1.0 and will be removed in 1.2.

If you wish to scale the data, use Pipeline with a StandardScaler in a preprocessing stage. To reproduce the previous behavior:

```
from sklearn.pipeline import make_pipeline
```

```
model = make_pipeline(StandardScaler(with_mean=False), Lasso())
```

If you wish to pass a `sample_weight` parameter, you need to pass it as a fit parameter to each step of the pipeline as follows:

```
kwargs = {s[0] + '__sample_weight': sample_weight for s in model.steps}
model.fit(X, y, **kwargs)
```

```
Set parameter alpha to: original_alpha * np.sqrt(n_samples).
warnings.warn(
```

Out[34]: Lasso(alpha=0.5, normalize=True)

```
In [35]: 1 coef = pd.Series(lasso_model.coef_, predictors).sort_values()
        2 print(coef)
```

```
fuel-type      -7268.421633
stroke         -4024.350234
bore           -2191.378272
engine-type    -648.687296
symboling      -603.255536
body-style     -294.850287
num-of-doors   -271.787484
city-mpg       -216.952356
make           -196.387898
wheel-base    -20.796517
length         11.544332
engine-size    135.243830
highway-mpg    260.129778
fuel-system    261.962291
height        319.023898
compression-ratio 525.214813
num-of-cylinders 579.079765
drive-wheels   614.931666
width         743.132522
aspiration     2861.847142
engine-location 12747.200906
dtype: float64
```

```
In [36]: 1 y_predict = lasso_model.predict(x_test)
        2 score = lasso_model.score(x_test,y_test)
        3 print('accuracy of lassoo model is ....',score)
        4 lasso_model_mse = mean_squared_error(y_predict,y_test)
        5 print('error of lasso model is...', math.sqrt(lasso_model_mse))
```

```
accuracy of lassoo model is .... 0.8549364877465662
error of lasso model is... 2997.4534137729383
```

**conclusion for lasso - accuracy is 85% and error is 2997 dollars**

```
In [38]: 1 from sklearn.linear_model import Ridge
2
3 ridge_model = Ridge(alpha = 0.5, normalize = True)
4 ridge_model.fit(x_train, y_train)
5 y_predict = ridge_model.predict(x_test)
6 score = ridge_model.score(x_test, y_test)
7 print('accuracy of ridge model is...', score)
8 ridge_model_mse = mean_squared_error(y_predict, y_test)
9 print('error of ridge model is ...',math.sqrt(ridge_model_mse))
10
```

```
accuracy of ridge model is... 0.9036806546176769
error of ridge model is ... 2442.4748979913443
```

C:\Users\91918\AppData\Local\Programs\Python\Python39\lib\site-packages\sklearn\linear\_model\\_base.py:141: FutureWarning: 'normalize' was deprecated in version 1.0 and will be removed in 1.2.

If you wish to scale the data, use Pipeline with a StandardScaler in a preprocessing stage. To reproduce the previous behavior:

```
from sklearn.pipeline import make_pipeline
```

```
model = make_pipeline(StandardScaler(with_mean=False), Ridge())
```

If you wish to pass a sample\_weight parameter, you need to pass it as a fit parameter to each step of the pipeline as follows:

```
kwargs = {s[0] + '__sample_weight': sample_weight for s in model.steps}
model.fit(X, y, **kwargs)
```

```
Set parameter alpha to: original_alpha * n_samples.
warnings.warn(
```

**conclusion for ridge algorithm - accuracy is 90% and error is 2440 dollars**

**final conclusion - i am recommending ridge algorithm for vehicle price prediction project**

```
In [ ]: 1
```